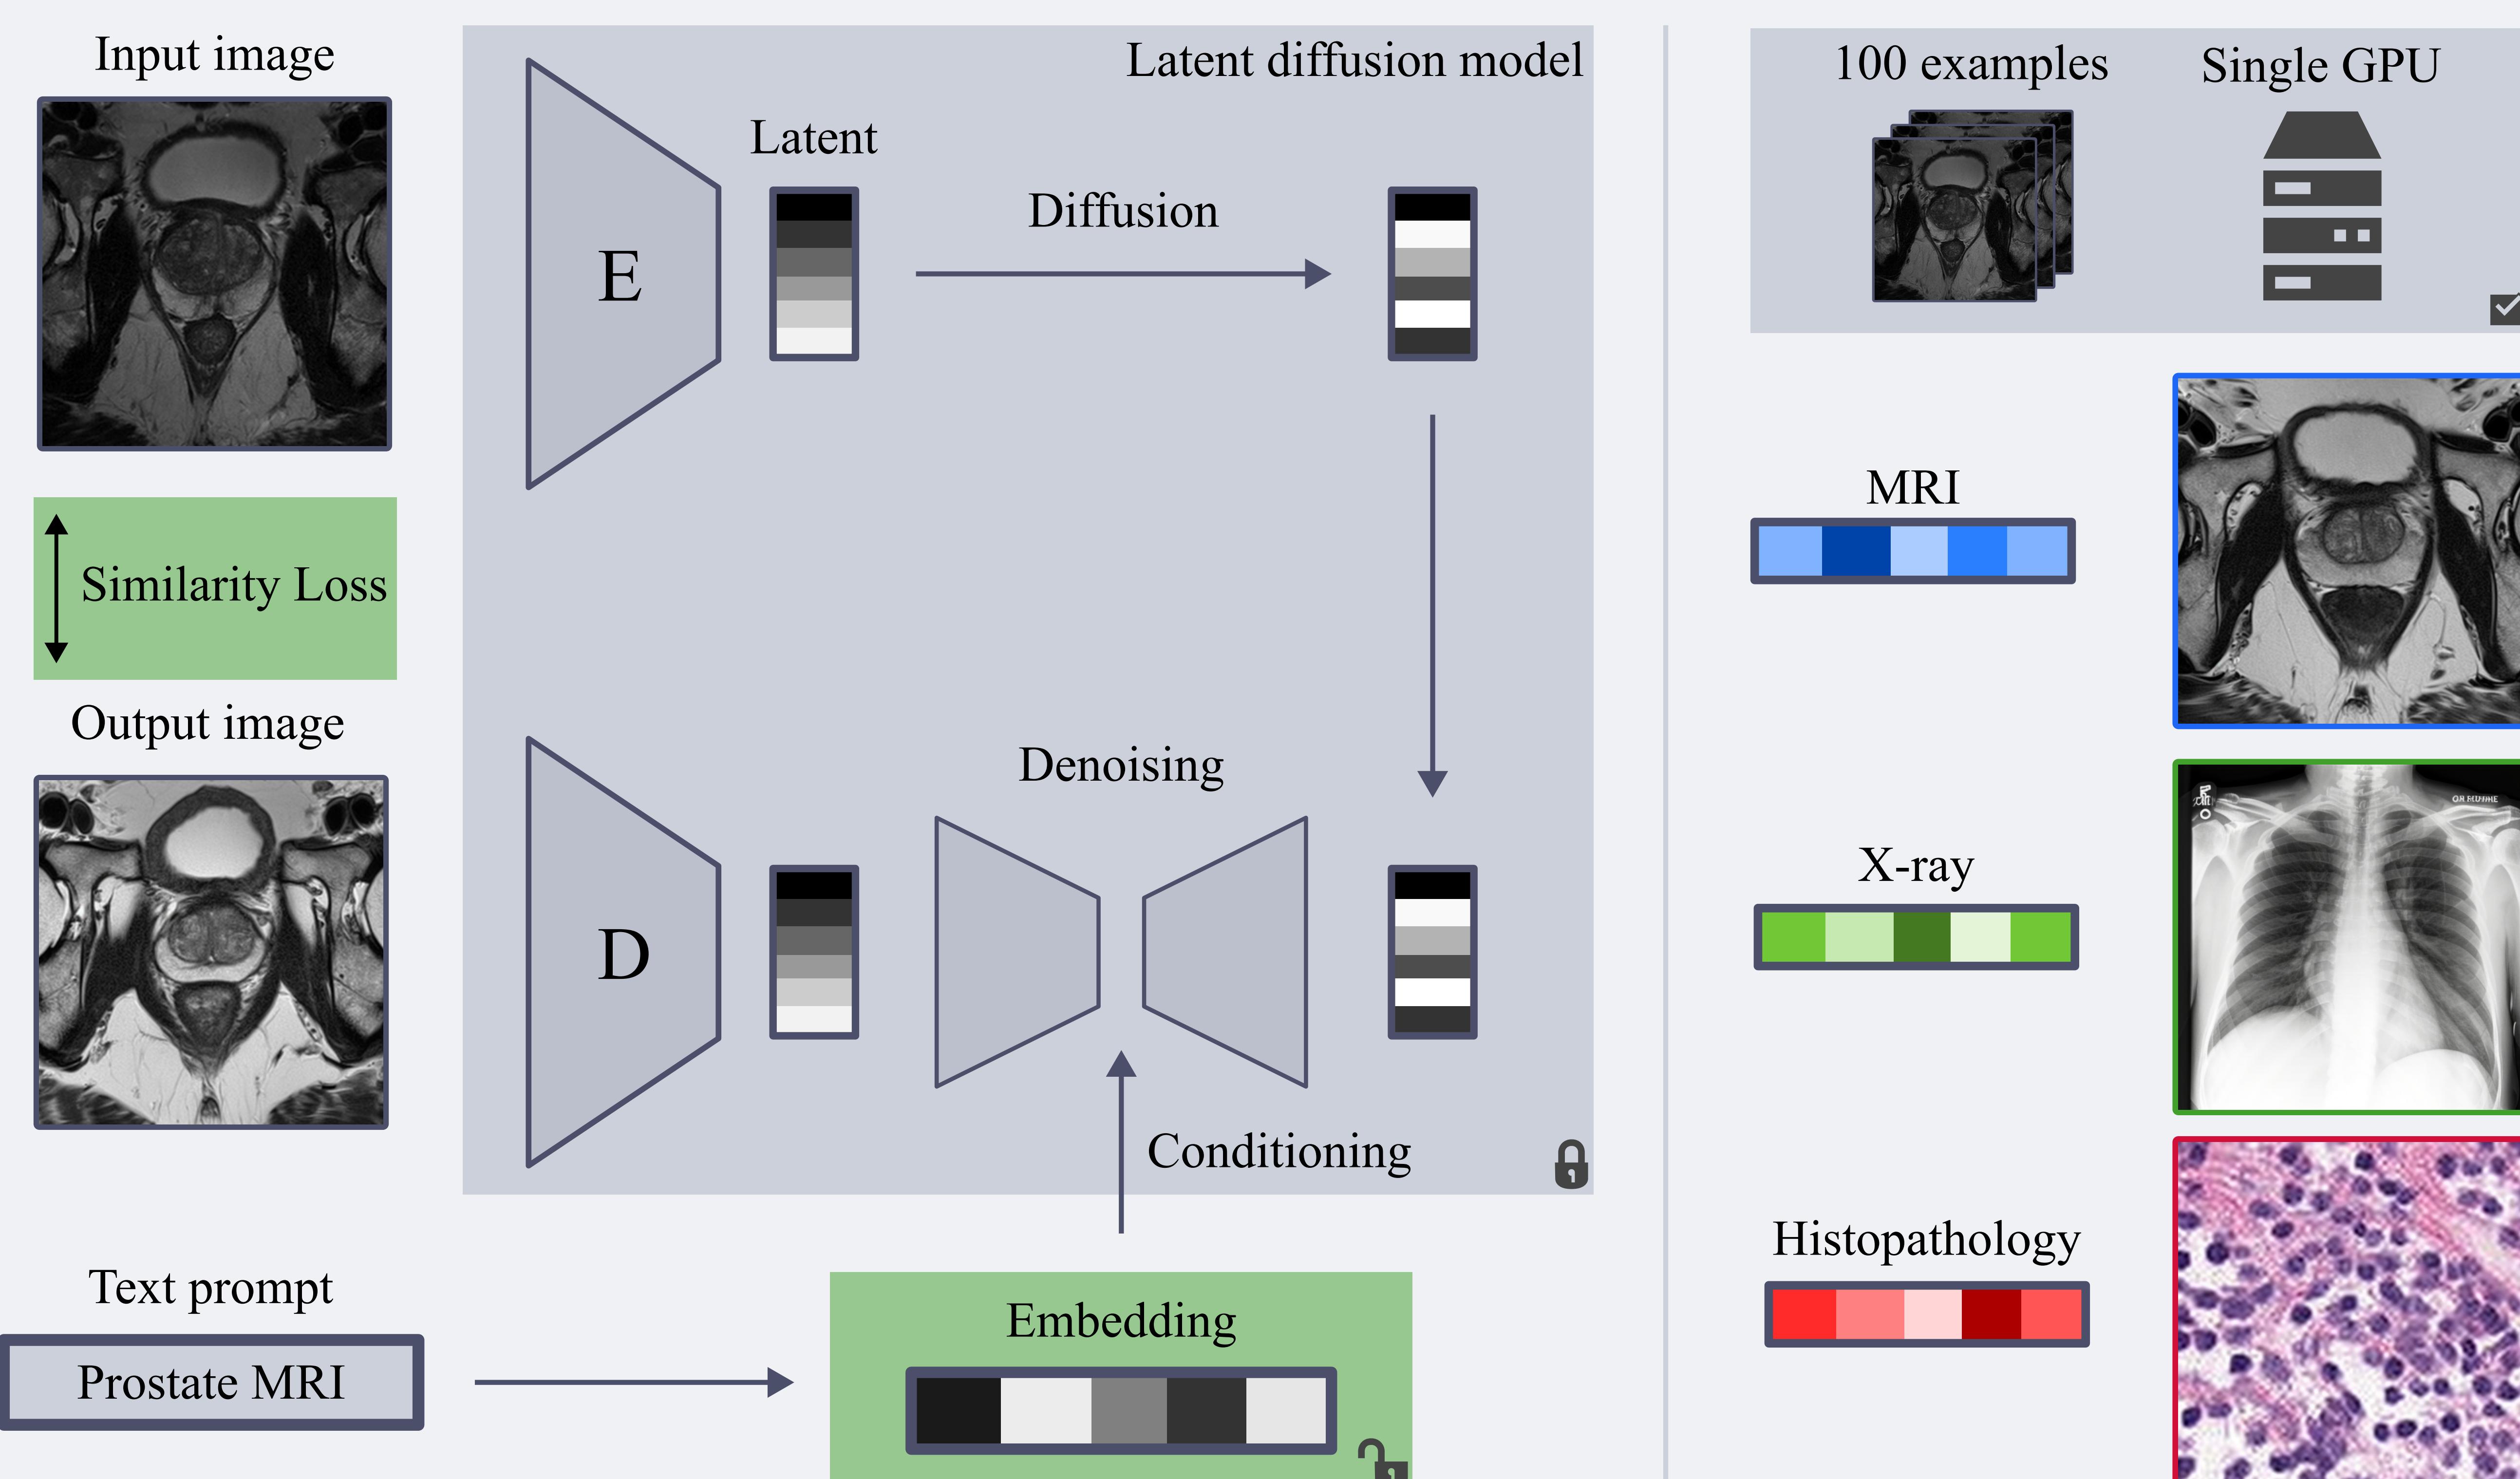


Generate faithful images in hours with 100 examples on 1 GPU

Medical diffusion on a budget: Textual Inversion for medical image generation

Bram de Wilde, Anindo Saha, Maarten de Rooij, Henkjan Huisman, Geert Litjens - Radboudumc, The Netherlands



INTRODUCTION

Diffusion models work great!

Inference possible on single GPU

Training requires **compute** and **data**

METHODS

Start from pre-trained Stable Diffusion

Adapt to medical modality with 100 images

Fine-tunes only the token embedding

RESULTS

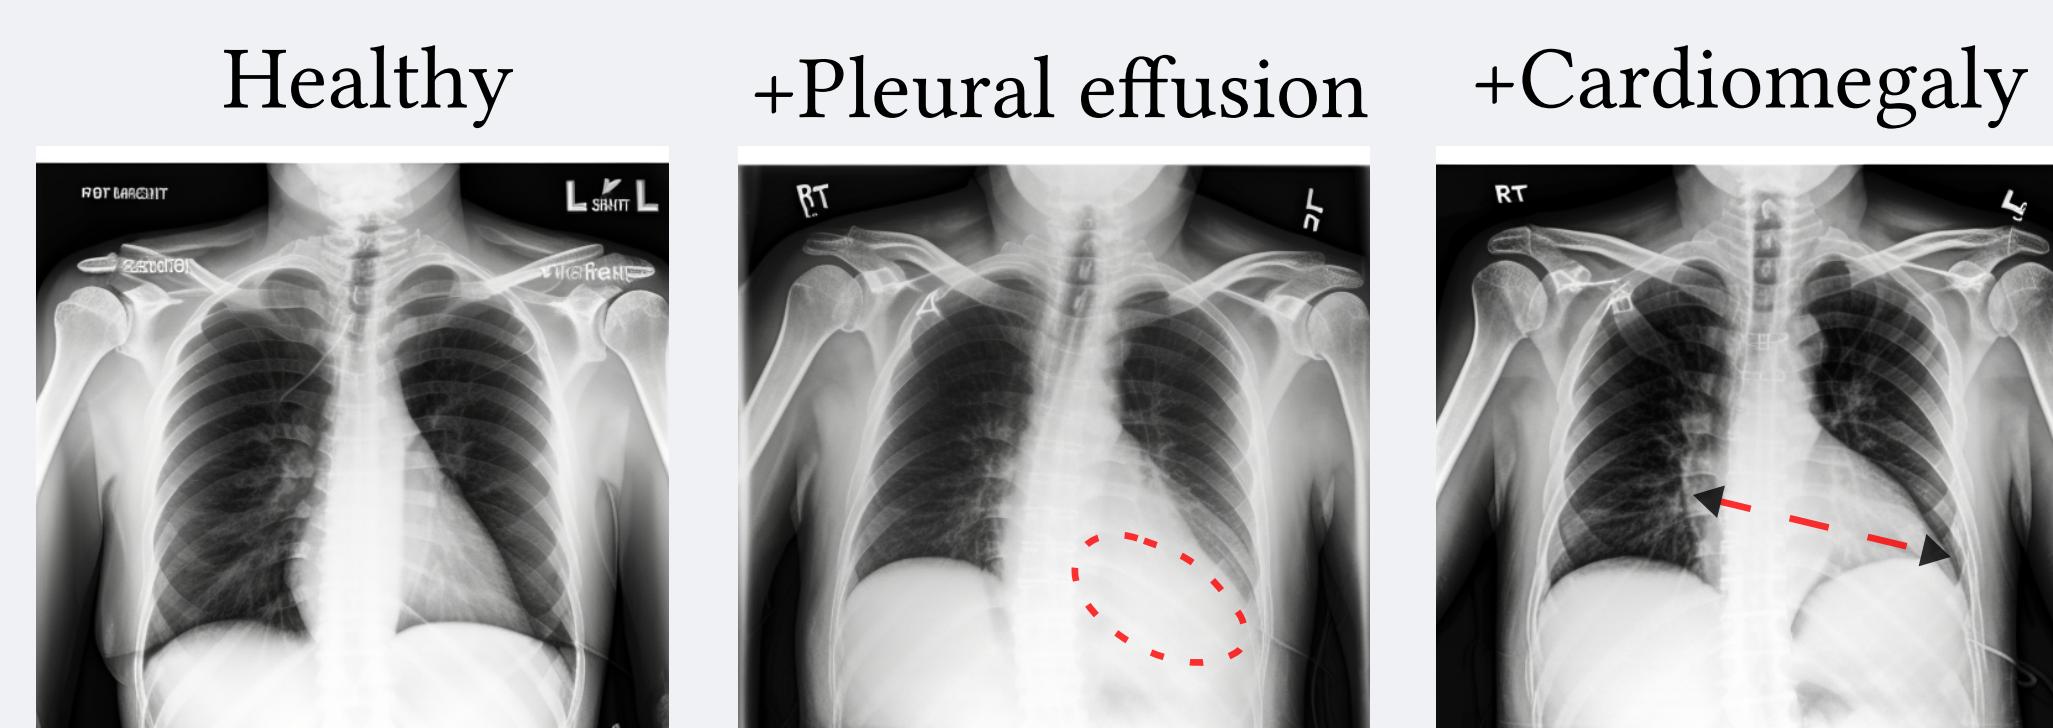
Use **larger embeddings** and **more images**

Images maintain or improve classification

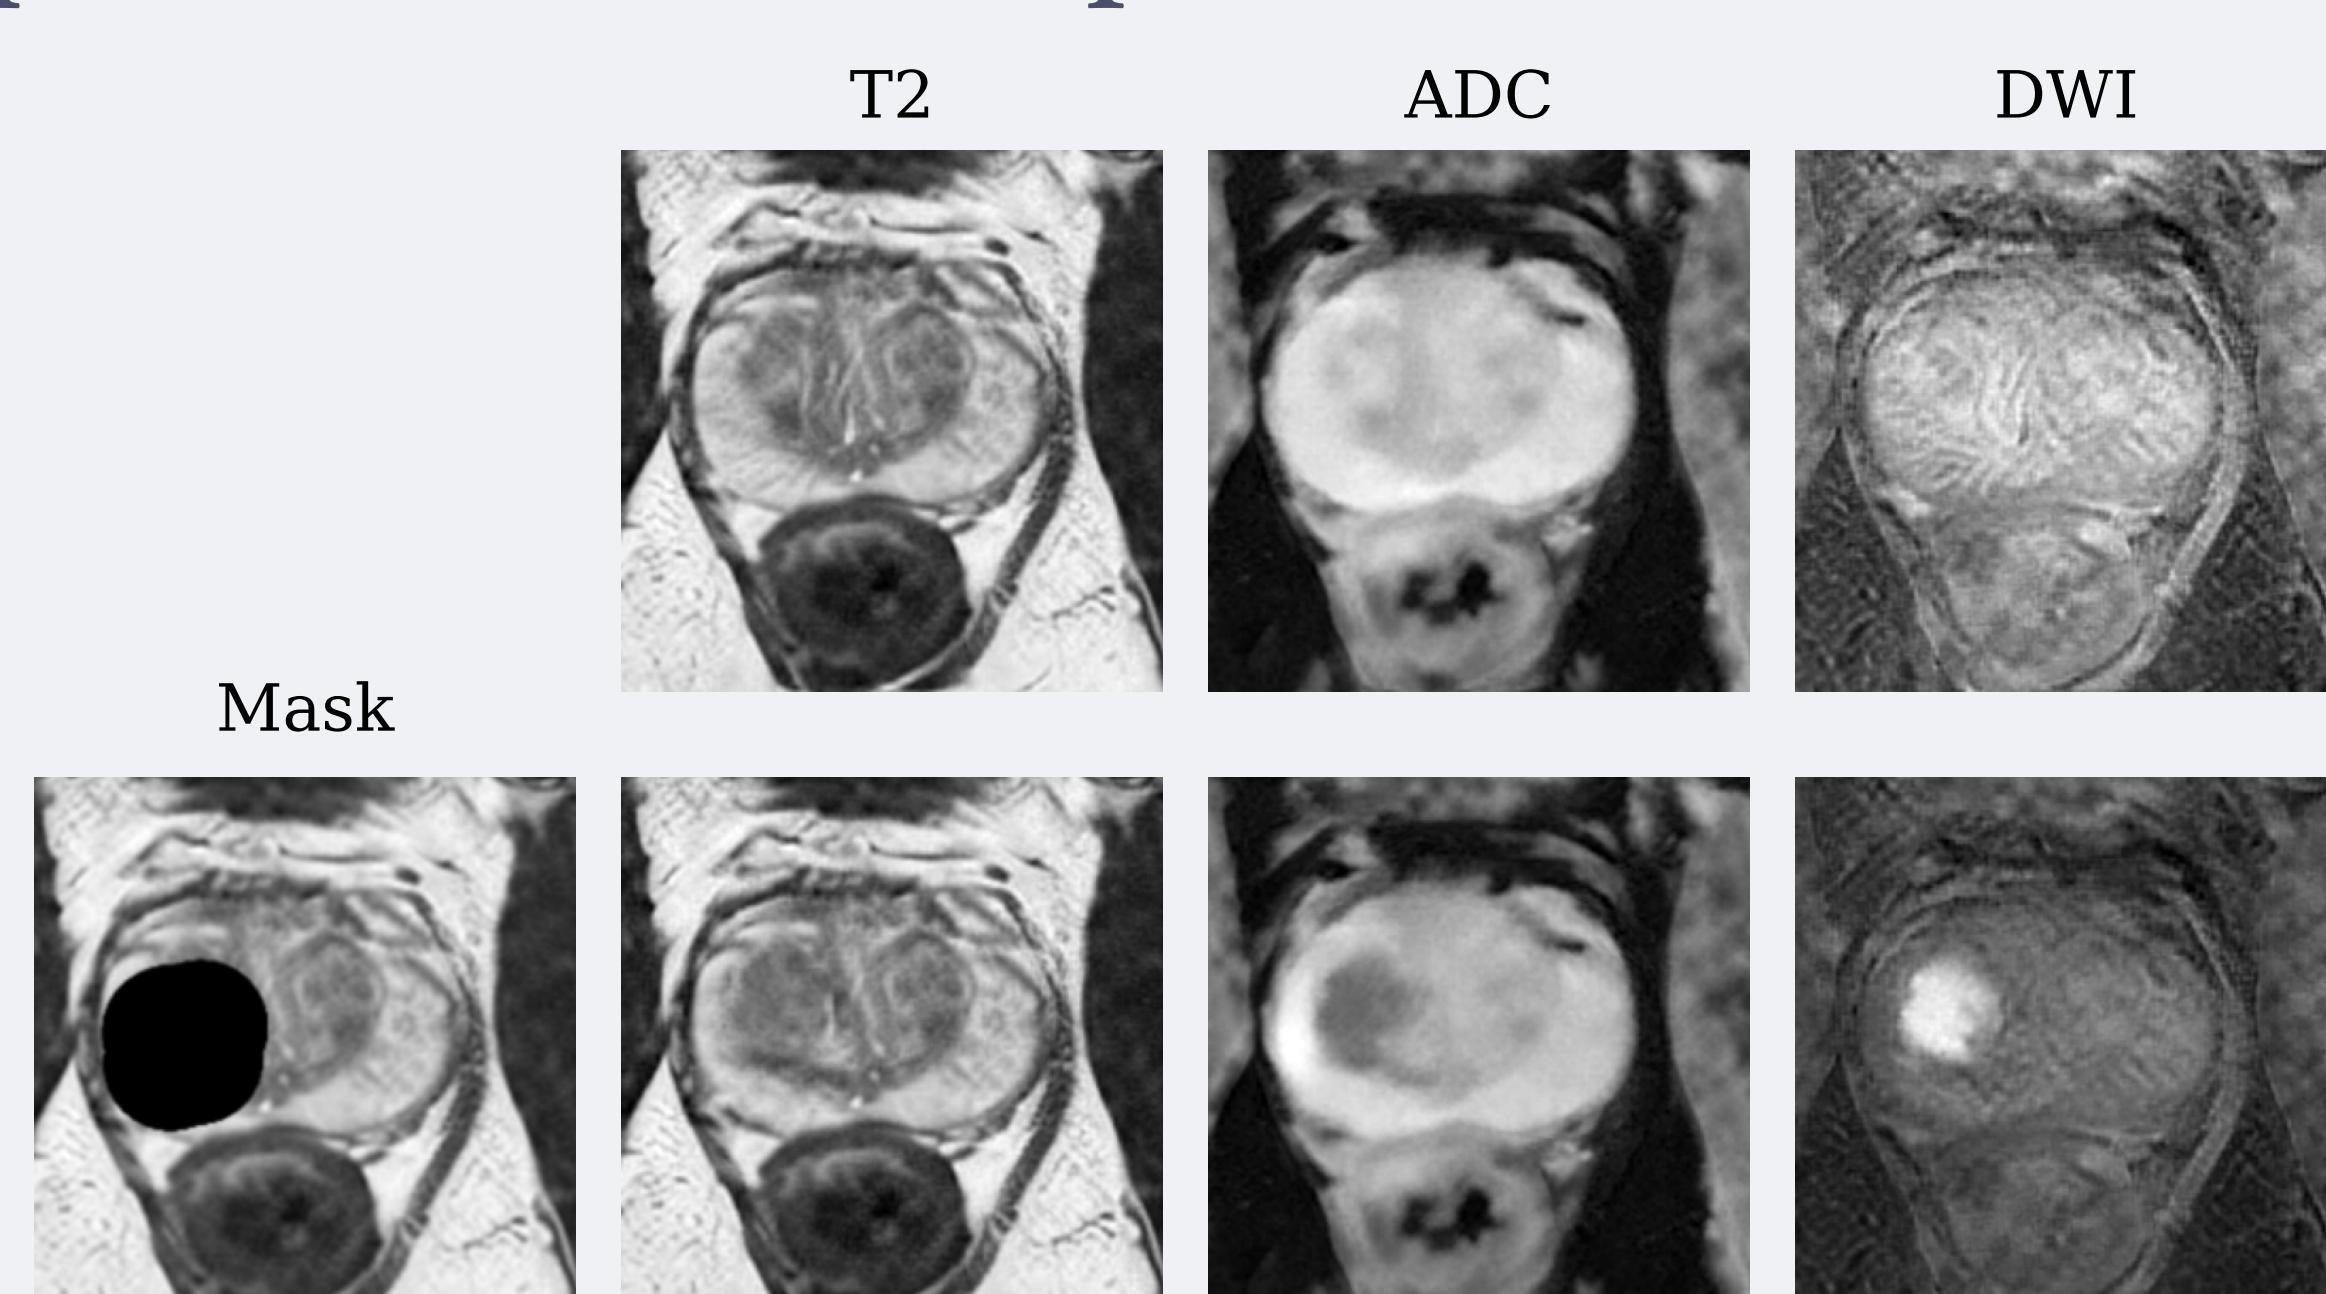
Images are **diagnostically accurate**

Preferred over StyleGAN3 by radiologist

Compose embeddings to combine diseases



Inpaint diseases at specific locations



← Scan for more details!

contact@bramdewilde.com